

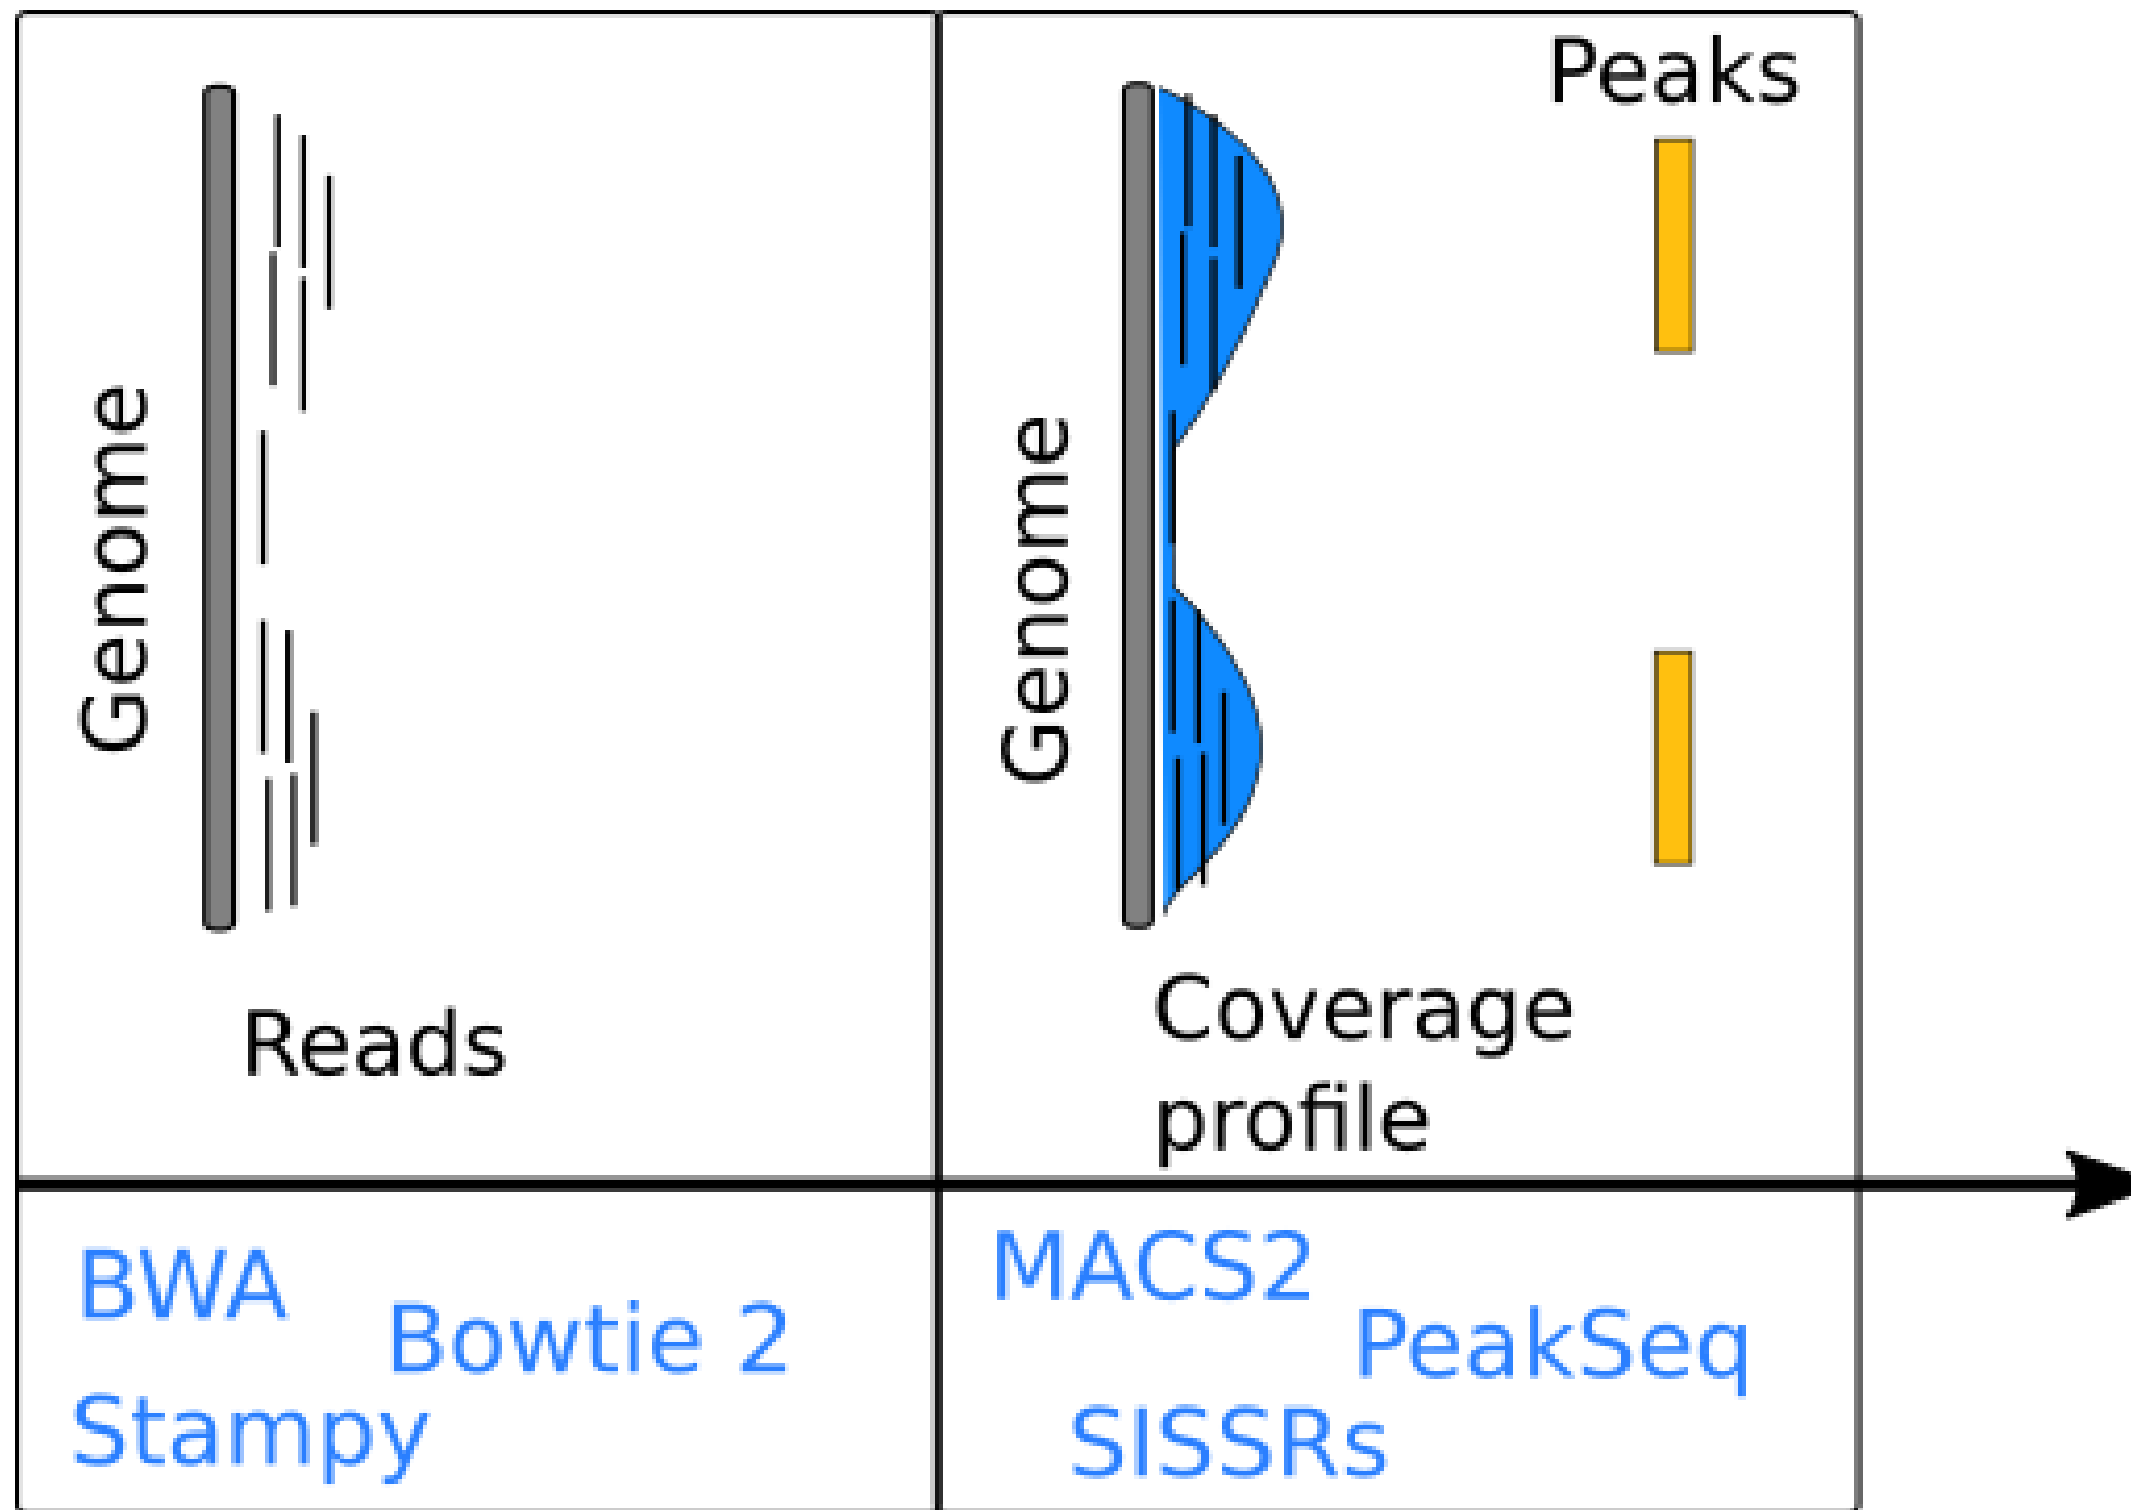
Importing data

CHIP-SEQ WITH BIOCONDUCTOR IN R



Peter Humburg

Statistician, Macquarie University



Handling sequence reads

- Usually stored in **B**inary **S**equences **A**lignment/**M**ap (BAM) format files.
- BAM record fields:
 - Read name: SRR1782620.7265769
 - Binary flag: 0
 - Reference sequence name and position of alignment: chr20 29803915
 - Mapping quality: 0
 - CIGAR string (alignment summary): 51M
 - Reference sequence and position of paired read (not used here): 0 0
 - Read sequence: AATGAAATGGAA ...
 - Read quality (ASCII encoded): CCCFFFFFFHHHH ...

Importing mapped reads into R

- Use `Rsamtools` package to interact with BAM files.
- `Rsamtools` provides functions for indexing, reading, filtering and writing of BAM files.

Use `readGAlignments` to import mapped reads.

```
library(GenomicAlignments)
reads <- readGAlignments(bam_file)
```

Returns `GAlignments` object.

Importing selected regions

- Use `BamViews` to define regions of interest.

```
library(GenomicRanges)
library(Rsamtools)
ranges <- GRanges(...)
views <- BamViews(bam_file, bamRanges=ranges)
```

- Then import reads as before.

```
reads <- readGAlignments(views)
```

The `BamViews` function supports multiple BAM files.

Importing peak calls

Use `import.bed` to load peak calls from a BED file.

```
library(rtracklayer)
peaks <- import.bed(peak_bed, genome="hg19")
```

Use `peaks` to define views into the BAM files.

```
bams <- BamViews(bam_file, bamRanges=peaks)
reads <- readGAlignments(bams)
```

Let's practice!

CHIP-SEQ WITH BIOCONDUCTOR IN R

Taking a closer look at peaks

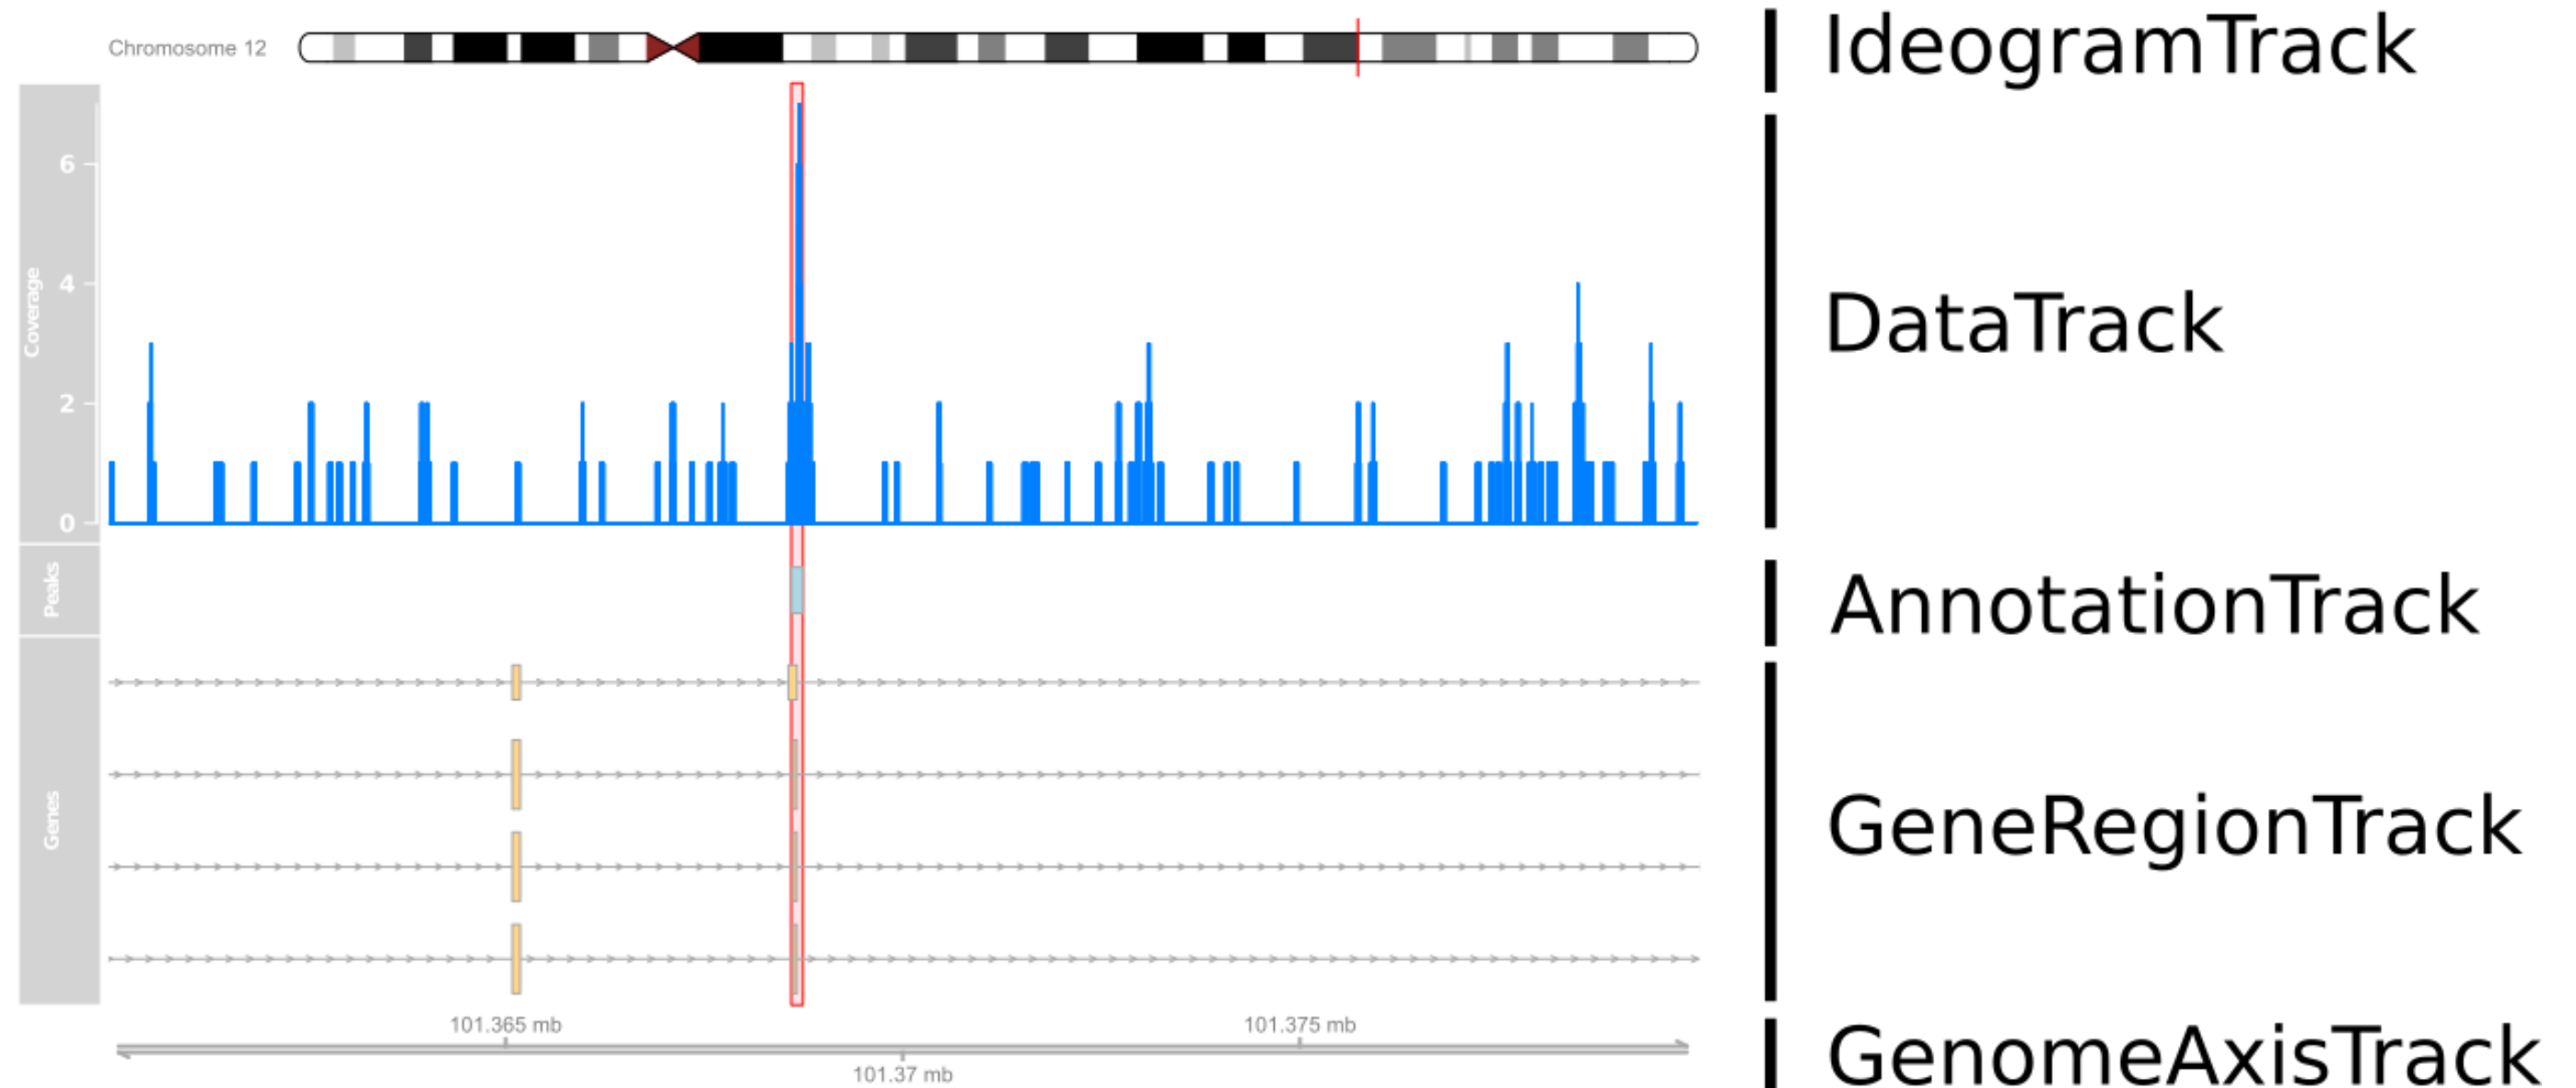
CHIP-SEQ WITH BIOCONDUCTOR IN R




Peter Humburg

Statistician, Macquarie University


Using Gviz




What are we plotting?



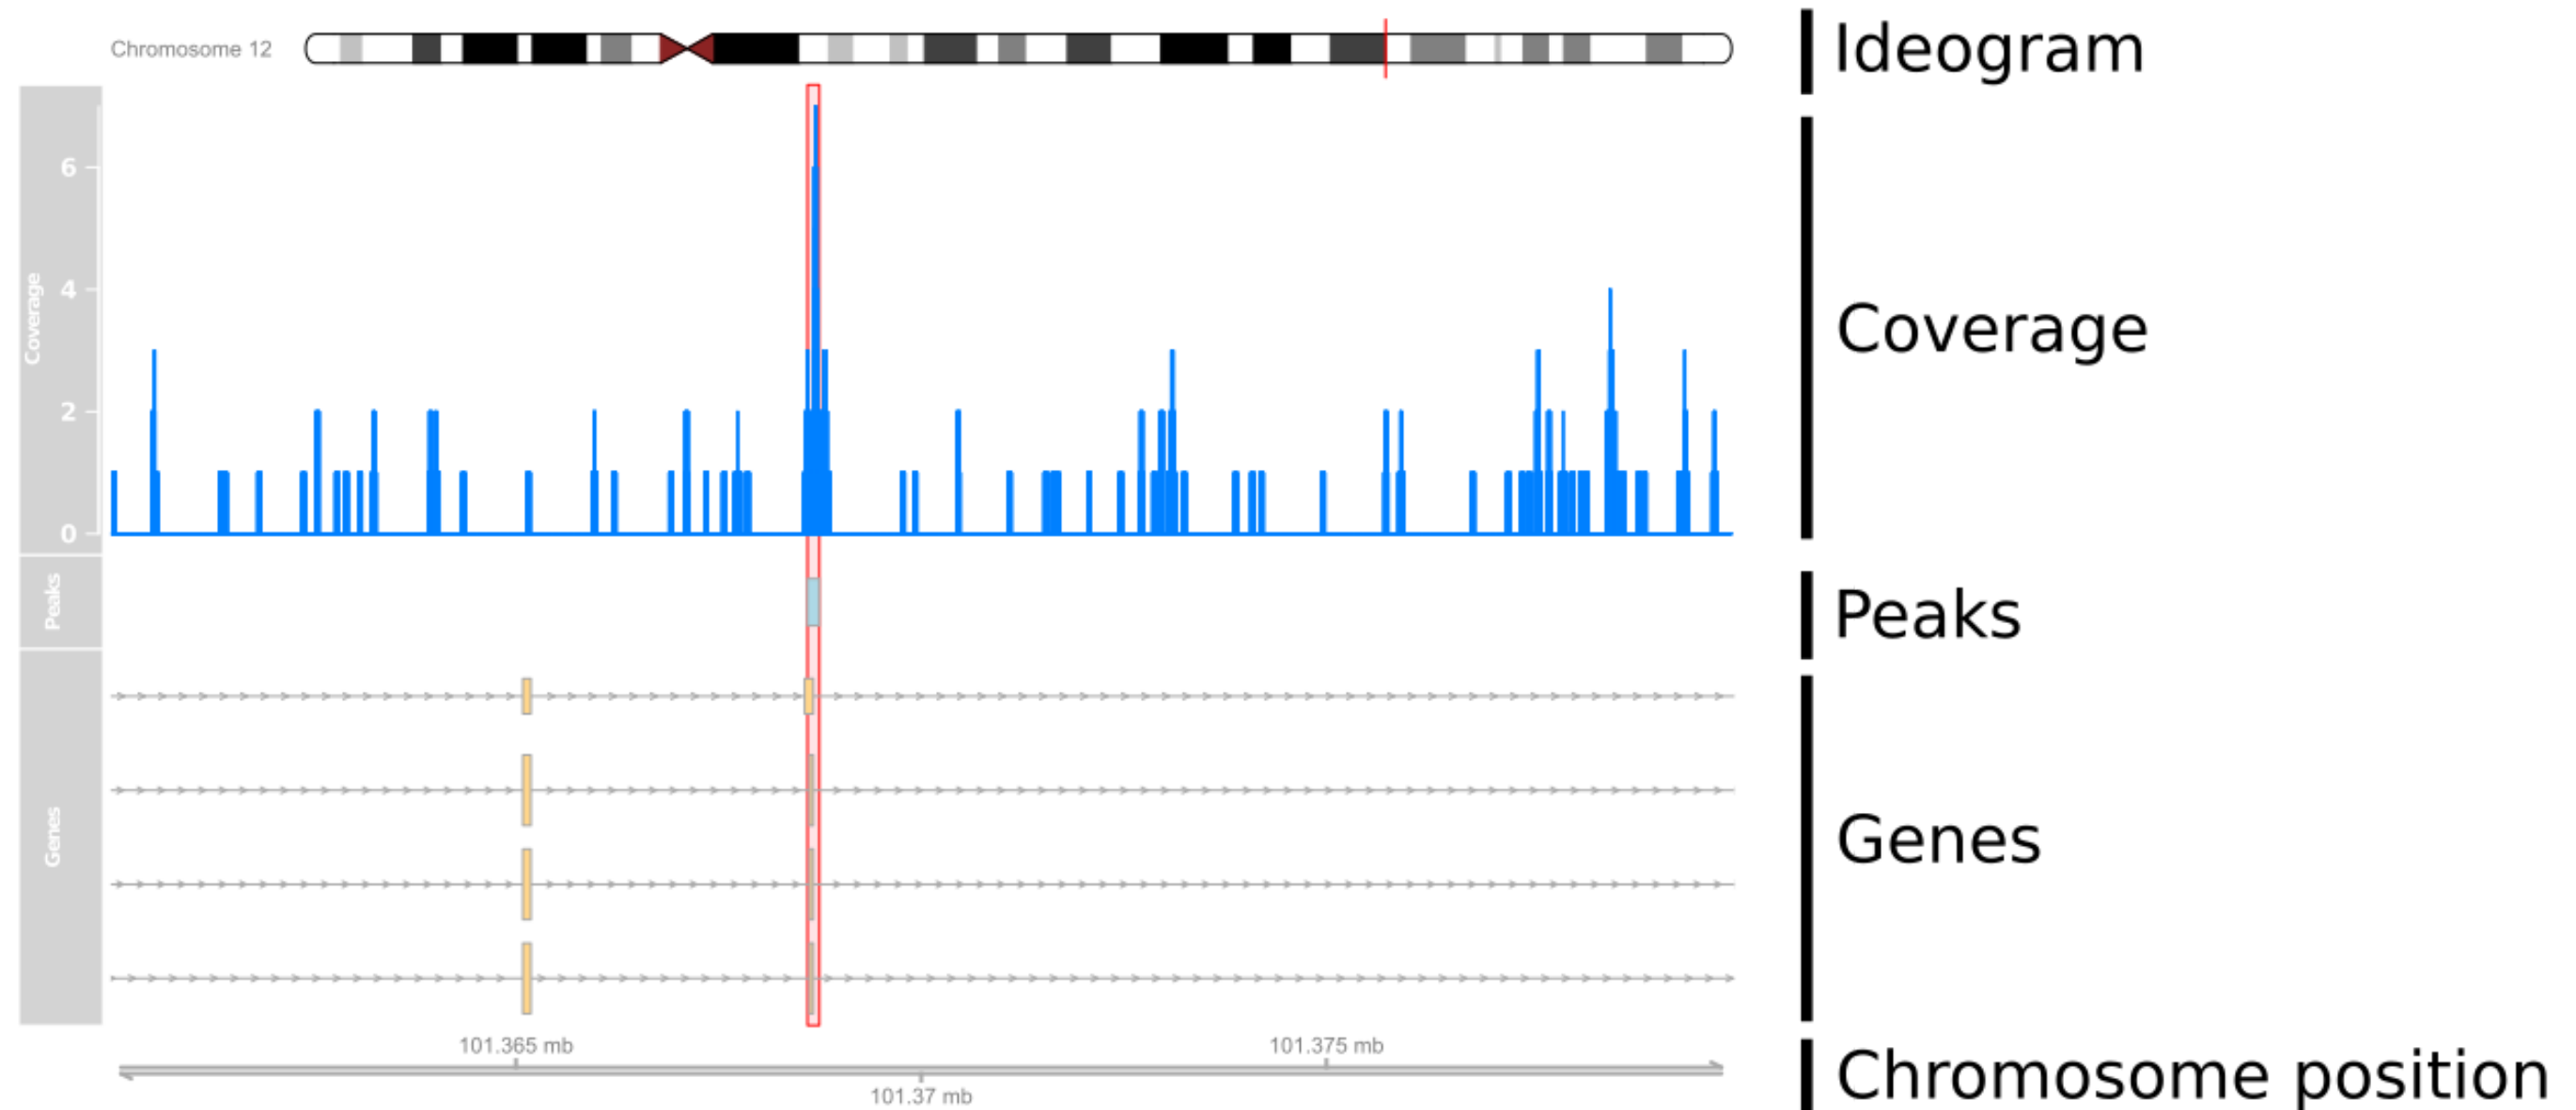
What are we plotting?



What are we plotting?




What are we plotting?



Setting-up coordinates

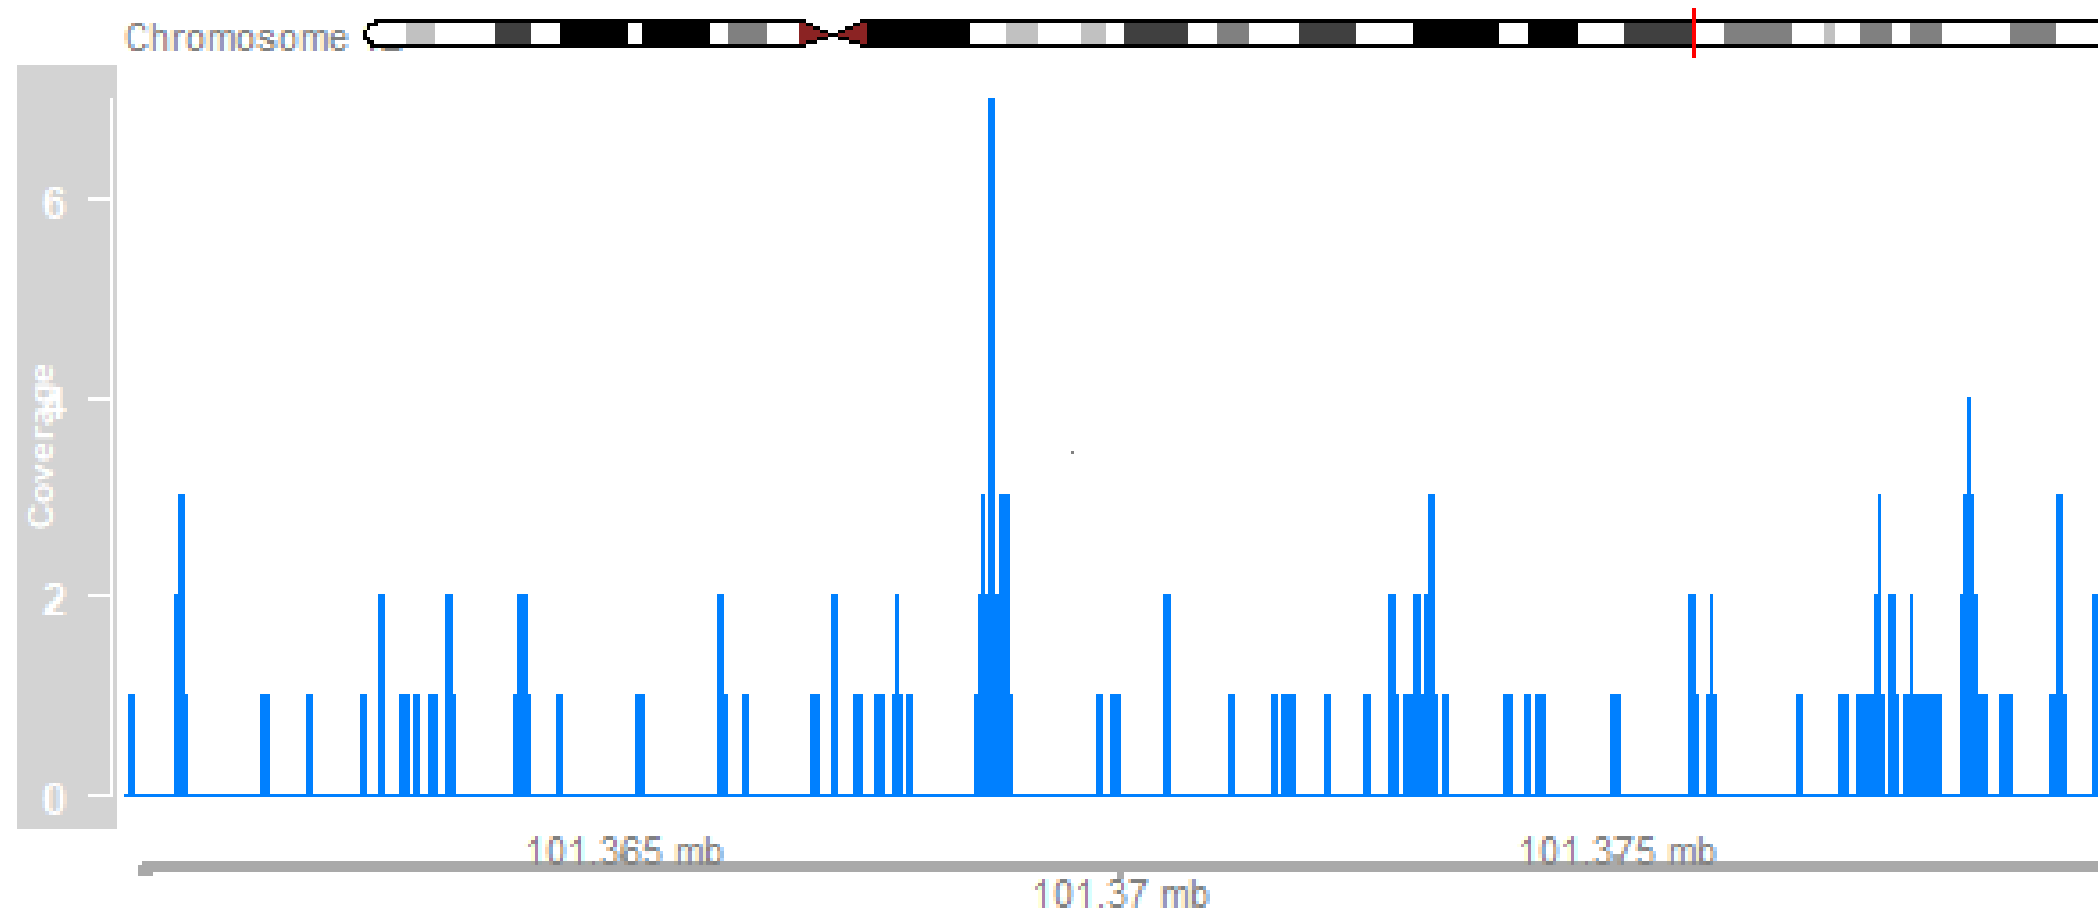
```
library(Gviz)
ideogram <- IdeogramTrack("chr12",
                          "hg19")

axis <- GenomeAxisTrack()
plotTracks(list(ideogram, axis),
           from=101360000,
           to=101380000)
```




Adding Data

```
cover_track <- DataTrack(cover_ranges,window=100000,type='h',name="Coverage")  
plotTracks(list(ideogram, cover_track, axis), from=101360000, to=101380000)
```




Adding Annotations

```
peak_track <- AnnotationTrack(peaks, name="Peaks")  
plotTracks(list(ideogram, cover_track, peak_track, axis),  
            from=101360000, to=101380000)
```



Gene Annotations

```
library(TxDb.Hsapiens.UCSC.hg19.knownGene)
tx <- GeneRegionTrack(TxDb.Hsapiens.UCSC.hg19.knownGene, chromosome="chr12",
                      start=101360000, end=101380000, name="Genes")
plotTracks(list(ideogram, cover_track, peak_track, tx, axis),
           from=101360000, to=101380000)
```



Let's practice!

CHIP-SEQ WITH BIOCONDUCTOR IN R

Cleaning ChIP-seq data

CHIP-SEQ WITH BIOCONDUCTOR IN R




Peter Humburg

Statistician, Macquarie University

Common Problems

Incorrectly mapped reads may produce false peaks.

- Genomic repeats.




Common Problems

Incorrectly mapped reads may produce false peaks.

- Genomic repeats.



- Incomplete reference sequence.




Common Problems

Incorrectly mapped reads may produce false peaks.

- Genomic repeats.




- Incomplete reference sequence.



- Low complexity regions.





Amplification Bias

- DNA fragments extracted from cells are copied multiple times prior to sequencing.
- Not all fragments produce the same number of copies.
- Multiple copies of the same fragment may be sequenced.
- A single DNA fragment may inflate coverage and lead to incorrect peak calls.

Quality Control Reports

```
library(ChIPQC)

qc_report <- ChIPQC(experiment="sample_info.csv", annotation="hg19")
ChIPQCreport(qc_report)
```

UP<>EXPAND ALLCOLLAPSE ALLSET AUTO WIDTHPRINT

ChIPQC Report

- Overview

This report was generated using [ChIPQC](#)

The report provides both general and ChIP-seq specific quality metrics and diagnostic graphics to allow for the quantitative assessment of ChIP-seq quality.

The report is split into three main sections:

- **QC Summary** - Overview of results.
- **QC Results** - Full QC results and figures.
- **QC files and versions** - Files and program versions used in QC

- QC Summary

Table 1. Summary of ChIP-seq filtering and quality metrics.

ID	Tissue	Factor	Condition	Replicate	Reads	Dup%	ReadL	FragL	RelCC	SSD	RiP%	RiBL%
GSM1598218	tissue: primary prostate tumor	AR	primary		12504249	2	51	183	0.11	0.47	0.1	5.1
GSM1598219	tissue: primary prostate tumor	AR	primary		16368472	1.9	51	185	0.1	0.57	0.042	6.1
GSM1598220	tissue: primary prostate tumor	AR	primary		11485090	1.7	51	220	0.66	0.5	4.6	5.3

Preparing input files

SampleID	Factor	Condition	Tissue	Treatment	bamReads	Peaks	PeakCaller
S1	AR	primary	primary prostate tumor	gleason score: 3+4=7	S1.bam	S1.bed	macs
S2	AR	primary	primary prostate tumor	gleason score: 3+4=7	S2.bam	S2.bed	macs
...

Cleaning the Data

- **Remove duplicate reads.**
- Remove reads with multiple hits.
- Remove reads with low mapping quality.
- Remove peaks in blacklisted regions.

Cleaning the Data

- Remove duplicate reads.
- **Remove reads with multiple hits.**
- **Remove reads with low mapping quality.**
- Remove peaks in blacklisted regions.

Cleaning the Data

- Remove duplicate reads.
- Remove reads with multiple hits.
- Remove reads with low mapping quality.
- **Remove peaks in blacklisted regions.**
 - Blacklisted regions are available from the ENCODE project.

Let's practice!

CHIP-SEQ WITH BIOCONDUCTOR IN R

Assessing enrichment

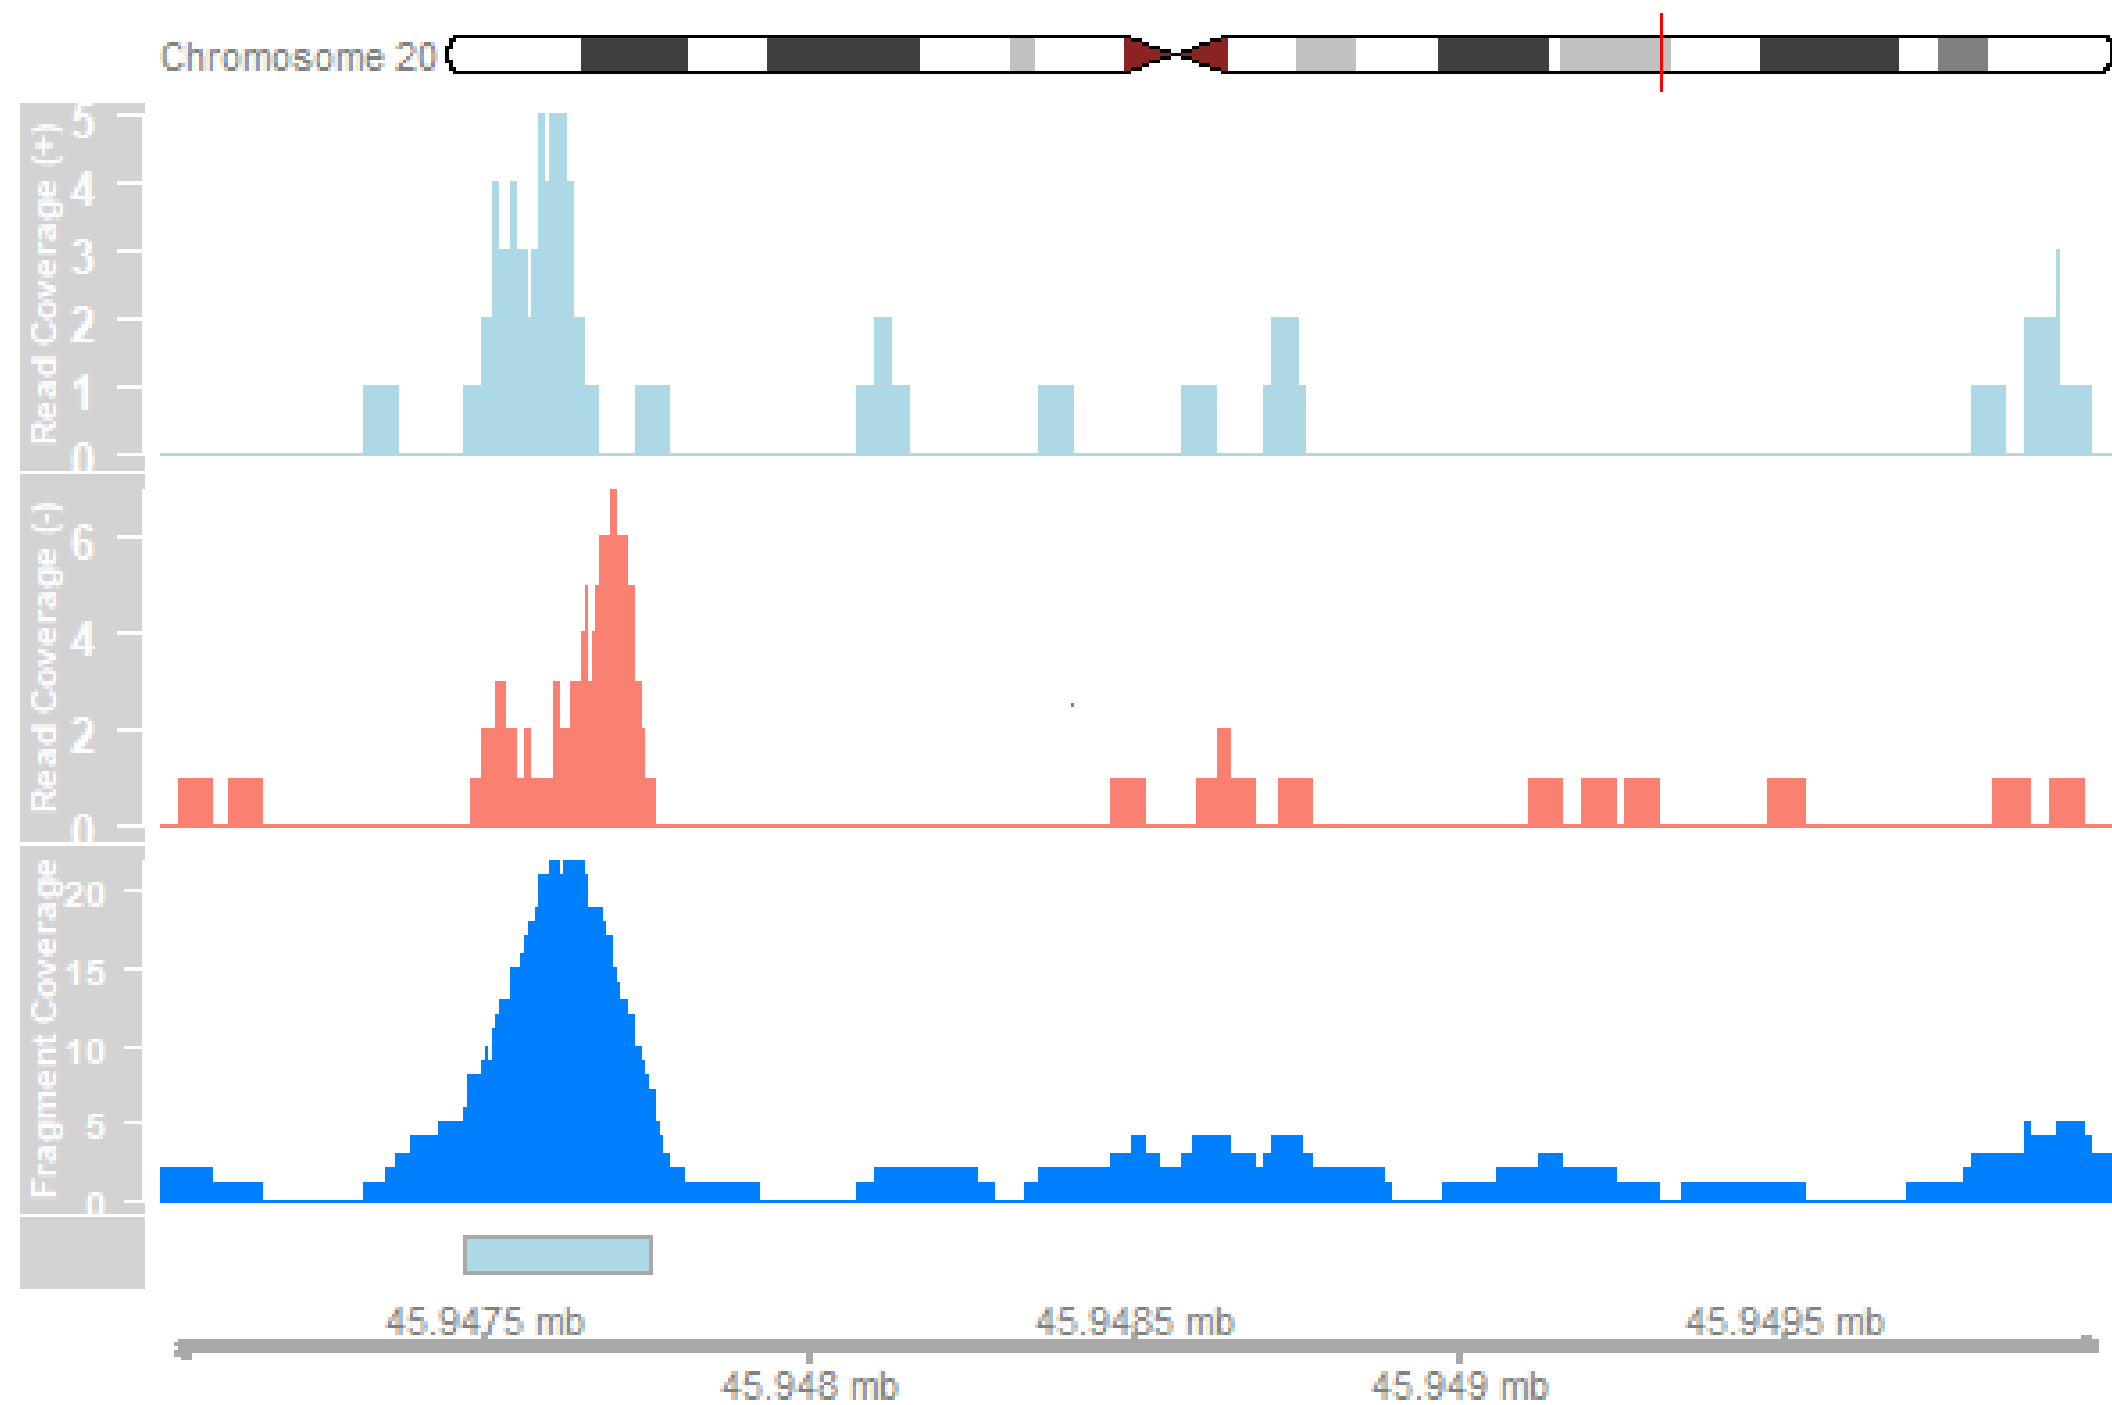
CHIP-SEQ WITH BIOCONDUCTOR IN R



Peter Humberg

Statistician, Macquarie University





Extending reads

Load the data:

```
reads <- readGAlignments(bam)
reads_gr <- granges(reads[[1]])
```

Obtain average fragment length:

```
frag_length <- fragmentlength(qc_report)["GSM1598218"]
```

Extend reads and compute coverage:

```
reads_ext <- resize(reads_gr, width=frag_length)
cover_ext <- coverage(reads_ext)
```

Coverage
profile


Genome

Peaks

Blacklist

Bins

Binned
data



Coverage for peaks

Create 200 bp bins along the genome.

```
bins <- tileGenome(seqinfo(reads), tilewidth=200,  
                  cut.last.tile.in.chrom=TRUE)
```

Find all bins overlapping peaks.

```
peak_bins_overlap <- findOverlaps(bins, peaks)  
peak_bins <- bins[from(peak_bins_overlap), ]
```

Count the number of reads overlapping each peak bin.

```
peak_bins$score <- countOverlaps(peak_bins, reads)
```

Binned coverage function

```
count_bins <- function(reads, target, bins){  
  # Find all bins overlapping peaks  
  overlap <- from(findOverlaps(bins, target))  
  target_bins <- bins[overlap, ]  
  
  # Count the number of reads overlapping each peak bin  
  target_bins$score <- countOverlaps(target_bins, reads)  
  target_bins  
}
```

Coverage for blacklisted regions

```
peak_bins <- count_bins(reads_ext, peaks, bins)
bl_bins <- count_bins(reads_ext, blacklist.hg19, bins)
```

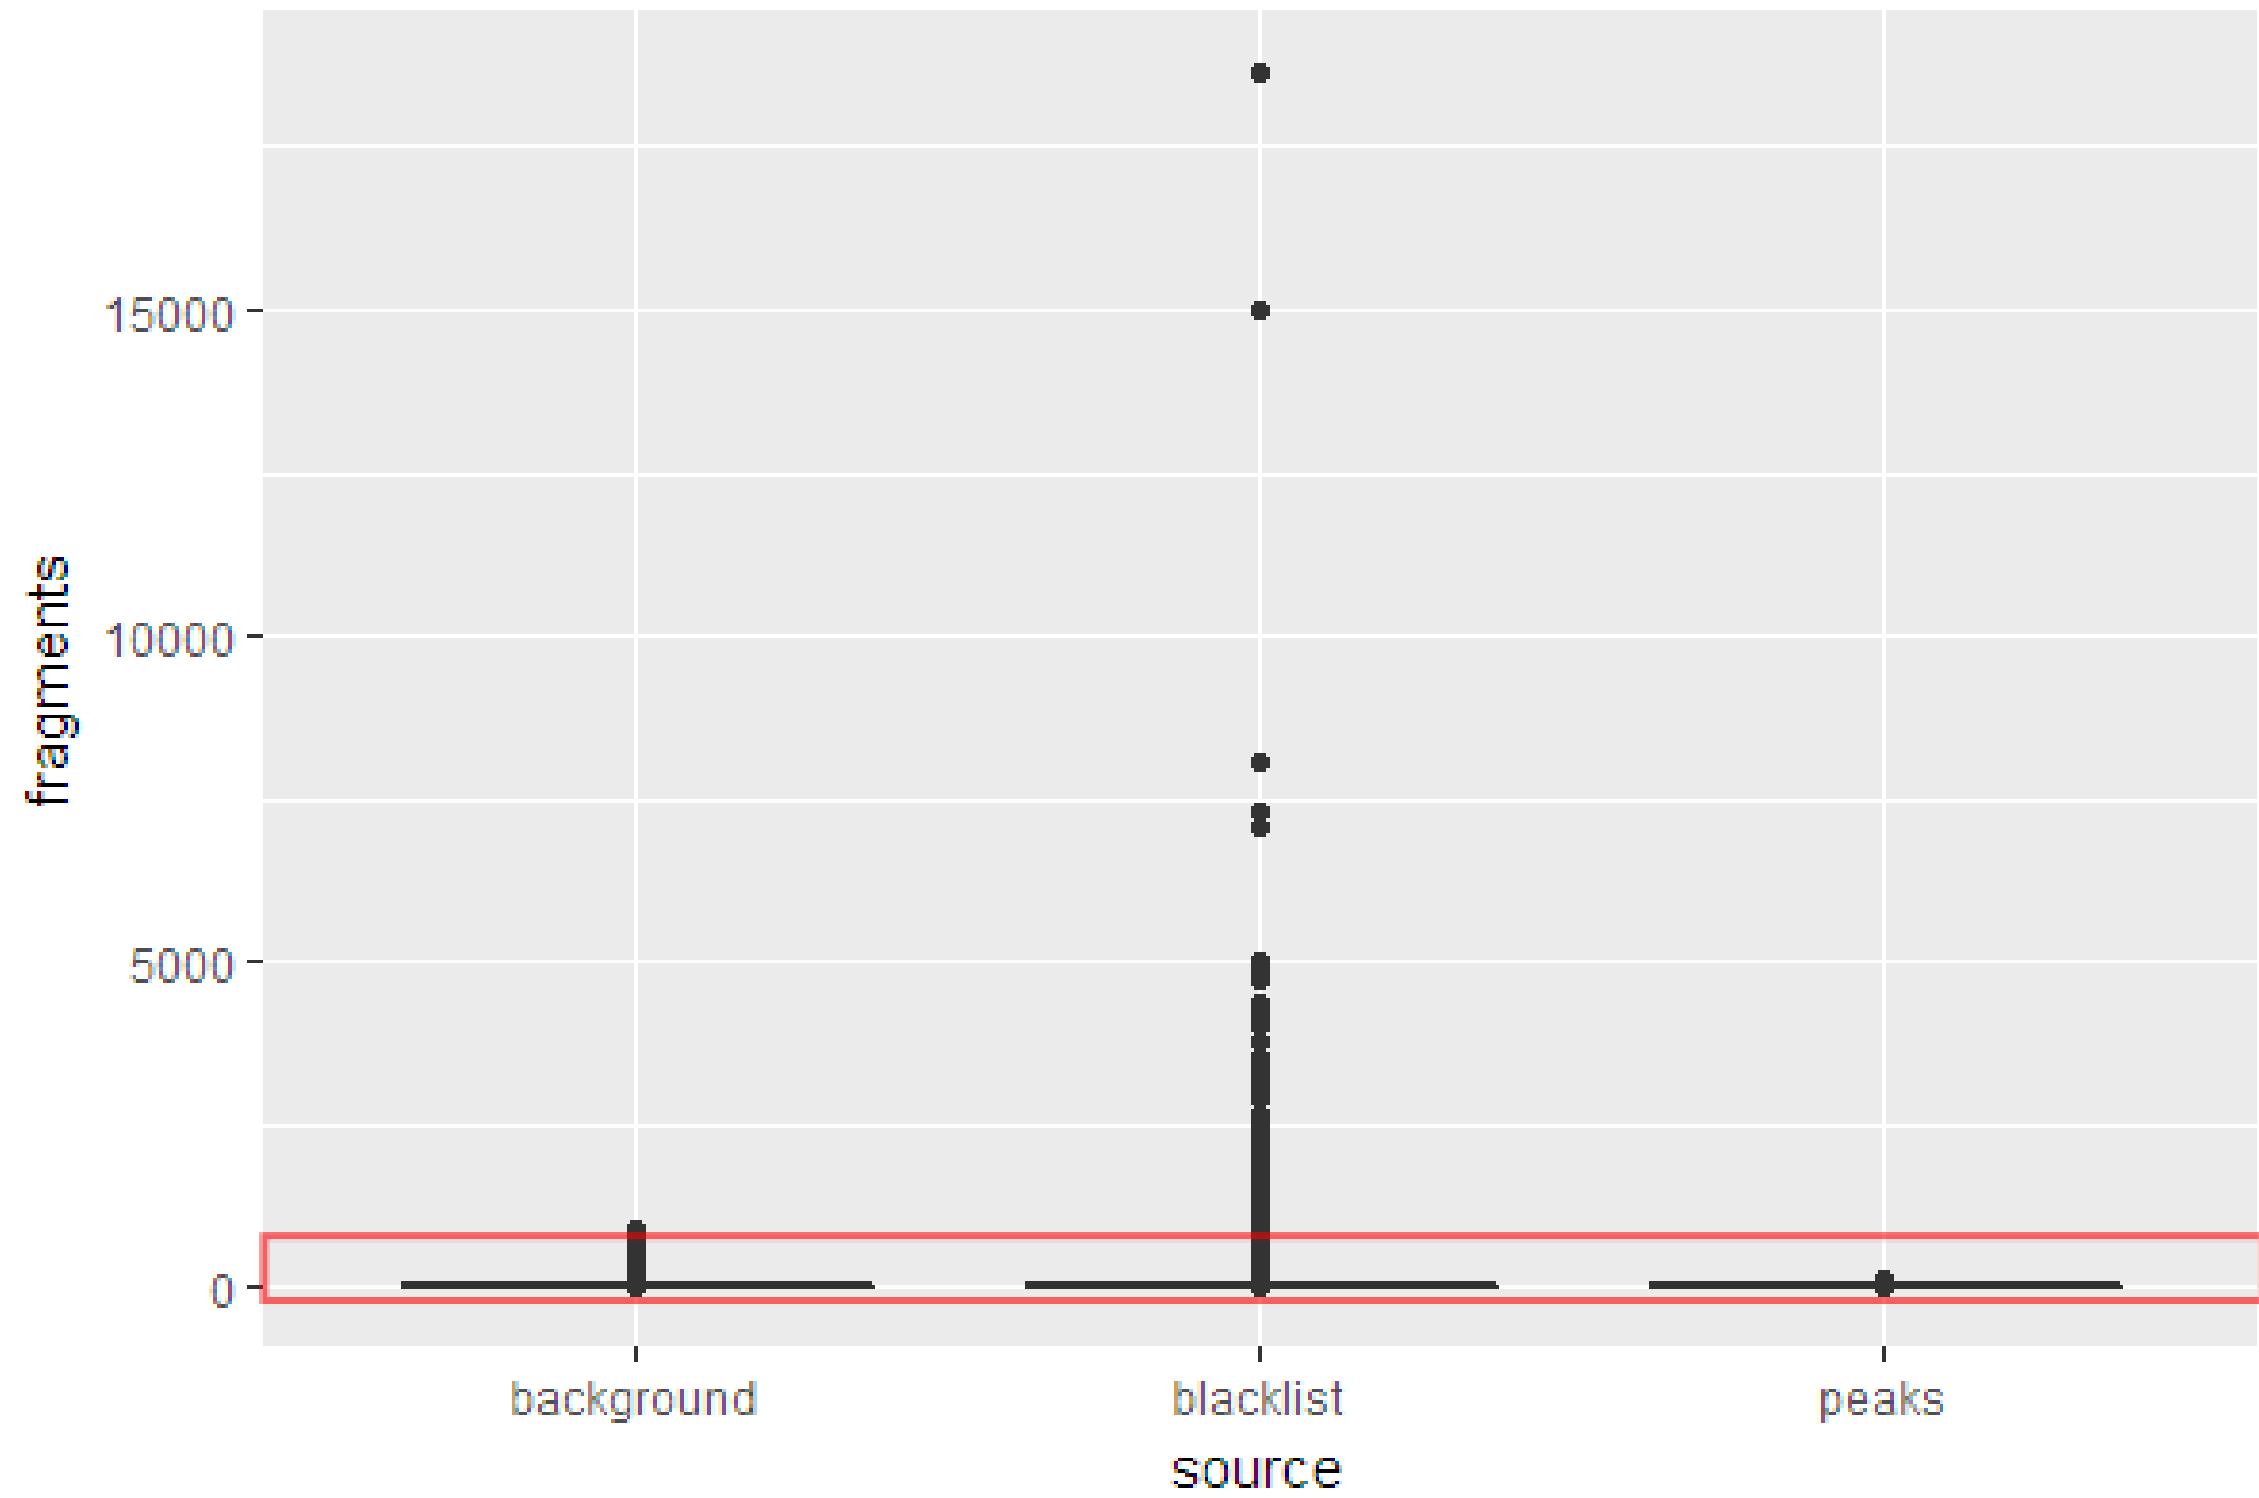
Background coverage

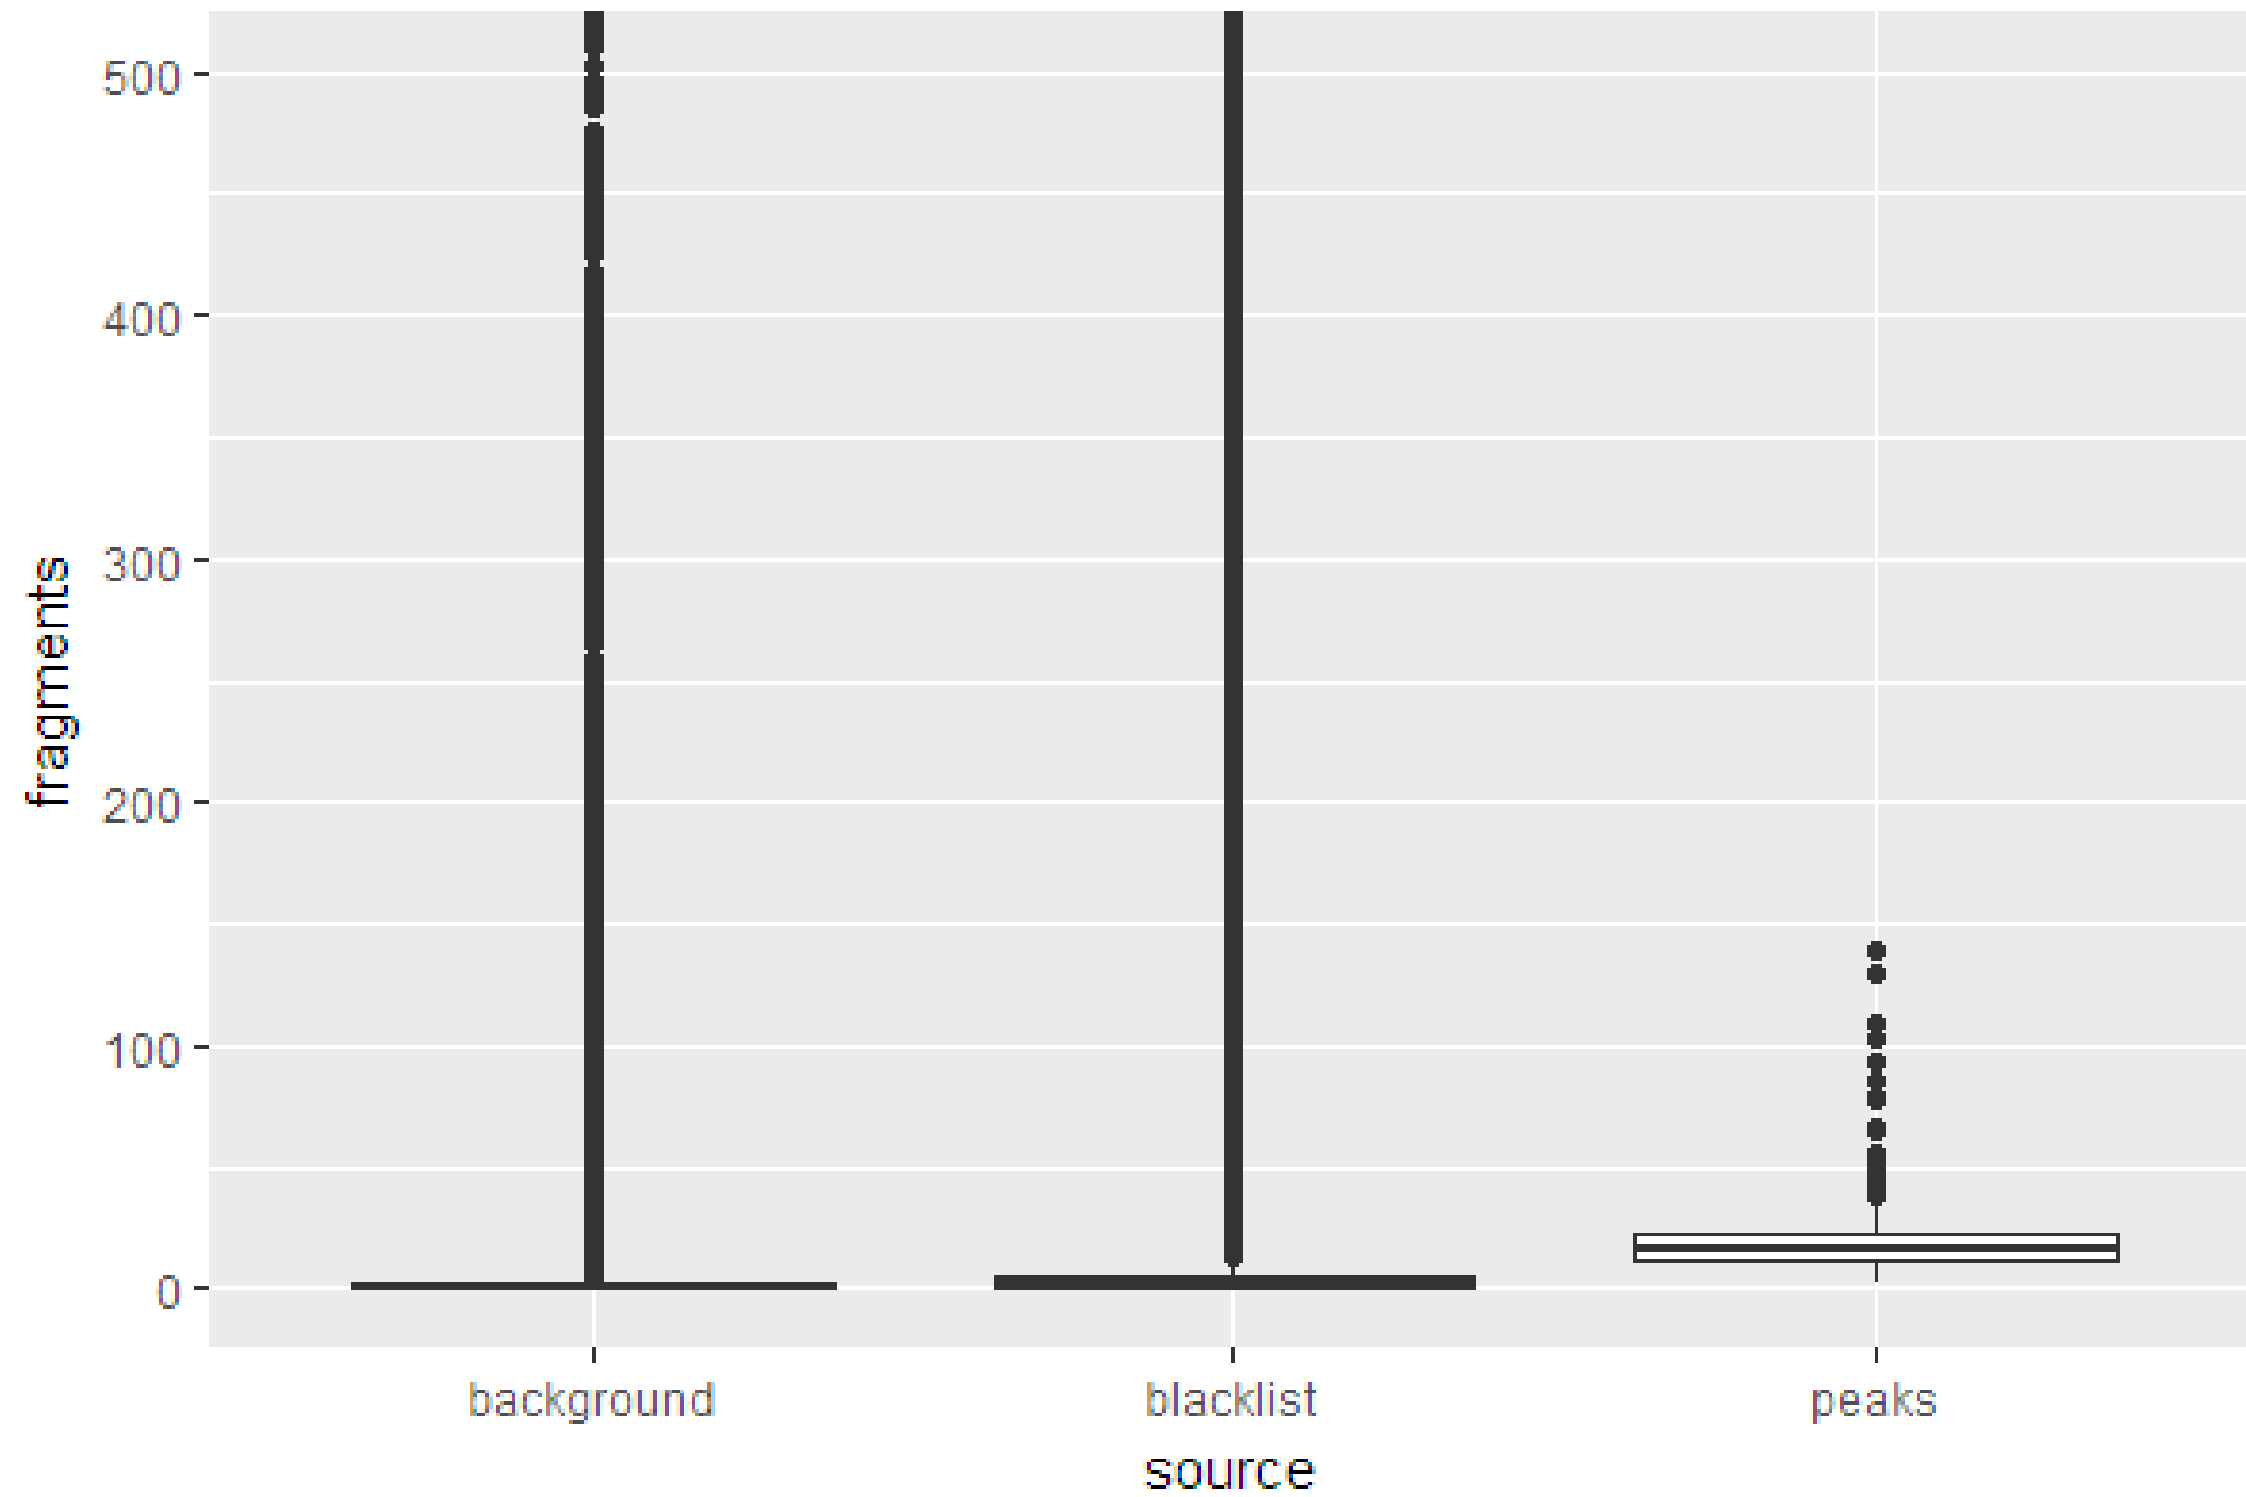
Remove all bins already accounted for.

```
bkg_bins <- subset(bins, !bins %in% peak_bins & !bins %in% bl_bins)
```

Count number of reads overlapping with each remaining bin.

```
bkg_bins$score <- countOverlaps(bkg_bins, reads_ext)
```



Let's practice!

CHIP-SEQ WITH BIOCONDUCTOR IN R